

Challenges and Solutions of Developing and Implementing a Novel Desktop-as-a-Service

Christian Baun and Johannes Bouché

Faculty of Computer Science and Engineering,
Frankfurt University of Applied Sciences,
Nibelungenplatz 1, 60318 Frankfurt am Main, Germany
[christianbaun|johannes.bouche]@fb2.fra-uas.de

Abstract. Desktop-as-a-Service (DaaS) is a cloud service category with tremendous potential for improving the everyday work of running and using applications. However, existing DaaS solutions have several drawbacks, which have limited the popularity of this service category in the past. Our novel DaaS solution, DESIGN, overcomes existing solutions' limitations and offers significant benefits. It can potentially revolutionize how we use applications, supporting non-modified Linux and Windows applications with just a browser on the client side. This paper describes some of the knowledge we gained during the development because of the challenges we faced and the solutions we found.

Keywords: DaaS · Compatibility · Performance · Stability · Usability

1 Introduction

This paper discusses challenges during the development and implementation of a novel Desktop-as-a-Service (DaaS) solution called DESIGN that enables the deployment and usage of unmodified Linux and Windows applications in the same way as web applications and can run inside public resources but can also be deployed in a private context. Since all interaction with DESIGN and the applications inside is done via a user's web browser, the users can use any client, no matter what hardware or host operating system it is compromised of. A browser is the only software component required to use the novel DaaS service.

Undoubtedly, developing a feature-rich DaaS system is a complex task, riddled with numerous obstacles and challenges. This paper meticulously categorizes these challenges into compatibility, performance, stability, and usability. It reveals these obstacles and presents the solutions we devised to overcome them, along with the valuable lessons we learned in the process.

This document is structured as follows. Section 2 discusses related work on DaaS solutions from an academic perspective, whereas section 3 describes the state-of-the-art DaaS architecture DESIGN we designed and implemented. Section 4 presents the most relevant challenges we faced, analyses our options for handling these, and describes the solution we found. Finally, section 5 discusses conclusions and includes directions for future work.

2 Related Work

DaaS has been a well-known service category since the emergence of cloud computing. Still, it got much less attention in research and literature. The challenges and obstacles when developing DaaS solutions, rather than infrastructure (IaaS) and platform services (PaaS), have seldom been discussed in the literature.

Celesti et al. [5] implemented in 2016 a DaaS using the IaaS solution OpenStack [21] and analyzed the characteristics and performance aspects of using noVNC, which is a client for the protocol Virtual Network Computing (VNC) that is implemented in the Hypertext Markup Language (HTML) and JavaScript, the protocol SPICE (Simple Protocol for Independent Computing Environments), and the remote desktop gateway Apache Guacamole as solutions for providing access to the desktop via a browser. One focus of the paper is the redirection of the sound interface of the virtual desktop. The paper's authors conclude that Guacamole, in combination with the Remote Desktop Protocol (RDP), is the best solution. The authors evaluated the lag time for remote audio playback in an Internet scenario by measuring the time between the automated initiation of the live streaming using the media player software VLC and the arrival of the music data by analyzing network traffic via the monitoring software Wireshark. When using Guacamole combined with the protocol RDP, the average lag was around 750 ms. When using Guacamole combined with the protocol VNC, the average lag was around 1750 ms. The paper also includes a short evaluation of the lag for video frames represented by background color changes of a Java application. Again, the configuration Guacamole and RDP offered the best performance video update with an average lag of around 300 ms. The configuration of Guacamole and VNC caused an average lag of around 500 ms.

Magana et al. [19] compared in 2019 the most popular remote desktop protocols and software implementations. The authors analyzed five protocols: PCoIP used in the Amazon WorkSpaces, RDP, TeamViewer, VNC, and Citrix Independent Computing Architecture (ICA). In the paper, the network transfer rate and its relation to the quality experienced by the DaaS user are evaluated by assuming three scenarios: using an office software suite, web browsing, and video streaming. One emphasis of the paper is comparing the transfer rate of the different protocols for three scenarios. In the Office Software Suite scenario, the transfer rate is measured and compared when performing several tasks like opening and editing a document (typing), loading an image, and saving a document. The web browsing scenario covers requesting three different web pages, and for the video streaming scenario, the same YouTube video file was viewed at different resolutions. The authors conclude that RDP requires fewer downstream and upstream network resources for all three scenarios. It is important to understand that the protocols RDP and VNC are implemented in open-source projects and proprietary solutions for many different operating systems. Depending on the various quality levels of the many existing implementations, it is difficult to generalize about the performance of the protocols, as individual implementations may have better or worse performance characteristics than expected.

Several works in literature have analyzed the latency of Cloud gaming services that have, in some aspects, similar requirements to DaaS offerings. According to Lampe et al. [17], the latency of locally executed games is compromised of the following components:

- Input lag caused by the controller device (e.g., mouse or keyboard).
- CPU time for processing the user input and the program itself.
- GPU time of the graphics processor for rendering the next frame.
- Time required for delivering the frame into the frame buffer.
- LCD response time (time required to display the frame on the screen).

In cloud gaming scenarios the following components increase latency further:

- Transfer time (upstream) for sending the user input to the service provider.
- Time required for capturing the frames and encoding it as a video stream.
- Transfer time (downstream) for sending the video stream to the client.
- Time required for decoding the video stream back into a frame.

For every deployment scenario or cloud service category possible, users expect a latency value that is not annoying when interacting with the service and its applications. Choy et al. [6], Claypool et al. [7], and Jarschel et al. [16] consider a maximum tolerable latency, based on subjective tests, being around 100 ms. This relatively low value sets the demand for a fast network interaction between client and service and limits their maximum distance. As mentioned in Clincy et al. [8], The former cloud gaming service provider OnLive specified in the year 2010 that the distance between the service provider and the consumer should not be greater than 1000 miles (approx. 1600 km) because otherwise, the round-trip communications delay time between will be too long for video games. A more in-depth comparison of commercially available cloud gaming services is given by Di Domenico et al. [10] or by Carrascosa et al. [4]. Additionally, implied delays and round-trips can be further reduced by using more sophisticated transmission protocols directly interacting with GPU drivers, such as Moonlight [14] [9].

On first thought, the latency requirements of a modern computer game are hardly the same as those of a desktop application. However, just as with computer games, the lagging operation of desktop programs has a negative impact on the user experience, and with video conferencing or other graphics-intensive applications, the utilization of resources and latency requirements are very similar to computer games.

3 Architecture

Figure 1 shows the architecture of our novel DaaS solution DESIGN that can run unmodified Linux and Windows applications in Linux containers (Linux and Windows applications compatible with Wine) and virtual machines (Windows applications incompatible with Wine).

The open-source server virtualization platform Proxmox VE runs bare-metal and offers virtual machines with the Kernel-based Virtual Machine (KVM) and

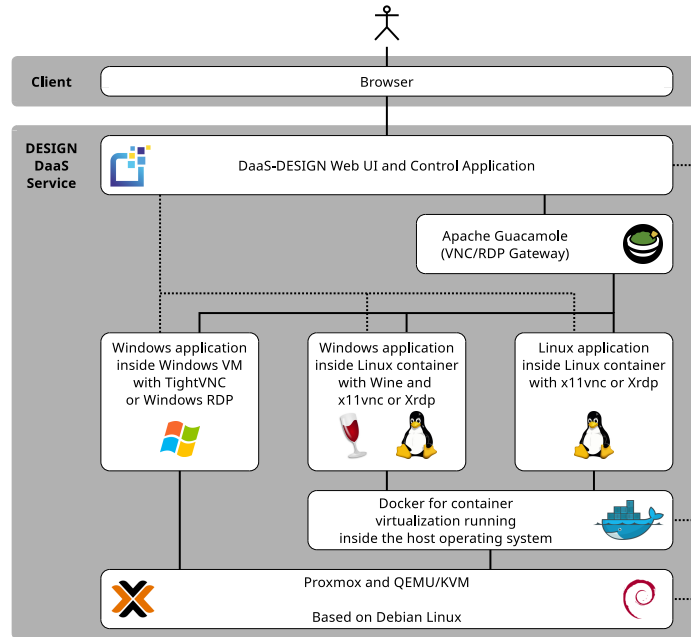


Fig. 1. Architecture of the novel DaaS DESIGN [2].

container-based virtualization via LXC Linux container. Since the LXC API is, in comparison to the Docker API, rather poorly documented and lacks several features, we considered Docker a better container virtualization solution. A deployment of docker directly in the host-operating system or inside a virtual machine is both possible.

Exporting the graphical user interface of Linux and Windows applications is possible using VNC or RDP protocols. Free server implementations exist for both operating system families. Since our DaaS solution focuses on exporting only the application’s graphical user interface and not full desktops – a feature not all server implementations include – the number of available projects and solutions is limited. For Linux containers, the free software projects x11vnc and Xrdp can be used in the DESIGN DaaS. For Windows VMs, the free software TightVNC and the Windows-internal RDP server can be used.

One of the main goals of DESIGN is to allow all interactions to be carried out solely using a browser. Thus, a broker or proxy software that mediates between the VNC or RDP implementation and the browser is required. The most advanced solution available is the free software project Apache Guacamole, which supports VNC and RDP and includes further relevant features regarding sound and printing.

The core components of DESIGN are our self-developed Web UI and the Control Application, which carry out the communication between all service components.

Starting a virtual machine or container instance automatically includes an SSH service. Our self-developed instance component waits until the SSH service is ready to satisfy requests. If applications are to be started or modifications are to be carried out, the instance component sends these requests to the virtual machine or container. Modification requests include resolution changes and fetching host status information.

The self-developed instance component generally allows users to carry out essential administration and usage tasks that the operating system does not offer via a generalized API.

4 Challenges and Solutions

While developing the novel DaaS platform DESIGN, we faced multiple obstacles and challenges. Each issue belongs to one of four categories: compatibility, performance, stability, and usability. This section presents the issues we faced, the possible approaches we analyzed, and the solutions we chose.

4.1 Compatibility

The first relevant characteristic of a DaaS that a user is confronted with is the list of compatible applications. In the best case, all applications of all popular operating systems can be integrated and used with such a novel platform. Several DaaS projects and solutions from the last two decades and similar concepts like web desktops or webtops lacked native support for popular applications.

Examples worth mentioning in this context include eyeOS and SilveOS. The project eyeOS [18, 24] simulates a desktop environment using HTML, the scripting language PHP, JavaScript, and the MySQL database management system. The project SilveOS [12] mimics the look and feel of a Windows desktop in the browser using the discontinued application framework Microsoft Silverlight. Despite being free software, both solutions have never become widely used since it is impossible to import and use native Windows or Linux applications there. All required applications have to be developed as web applications. This limitation caused most potential users to lose interest in web desktops because many popular native applications can hardly be replaced or re-developed, and many feature-rich and well-established Software-as-a-Service (SaaS) offerings like Microsoft Office 365 and Google Docs already exist.

A new DaaS is only accepted by a broad user group if it enables the integration of native applications in an unchanged manner. It is also desirable that a DaaS supports Linux and Windows applications both to allow flexible working and make the operation of a virtualization solution on the desktops unnecessary. Ideally, it should also be possible to run native Mac OS applications. However, this goal is difficult to achieve under legal conditions due to the Mac OS operating system's licensing restrictions.

Further essential features for most users are printing, sound, and compatibility with external storage media.

Guacamole supports redirecting audio, printing, and disk access via the protocol RDP. The printing feature does not allow access to local or remote printers; instead, it allows users to print directly to PDF files, which are received by the user’s web browser. As a prerequisite, the GhostScript interpreter must be deployed on the Guacamole server. File transfer over RDP supports Guacamole by emulating a virtual disk drive that persistently stores data on the Guacamole server. [13]

Compared to RDP, the protocol VNC offers fewer features in Guacamole since it does not support sound on its own. However, Guacamole supports sound using PulseAudio¹ as a workaround solution. VNC also does not offer file transfer and access to external storage media. Still, as a workaround solution, Guacamole supports file transfer via the SSH File Transfer Protocol (SFTP), which is an extension of the Secure Shell protocol (SSH).

4.2 Performance

Performance, from our viewpoint, stands out as the most crucial system property, implying the influence of other relevant properties, such as user experience, system responsiveness, and overall efficiency. Consequently, we encountered multifaceted challenges that potentially intersected across several layers within our architecture and which will be described in the following subsections.

Infrastructure and Platform In general, it can be said that the chosen infrastructure, platform, and utilized software combined may have the most impact on the overall resource consumption and are therefore of major importance. As a DaaS system combines all aspects of IaaS-, PaaS, and SaaS-based computation, each performance characteristic must be considered.

Regarding Infrastructure-as-a-Service (IaaS), our solution generally offers the flexibility to use either virtual machines or containers. Scientific publications discussing relevant performance aspects of such systems highlight the apparent advantage of a lightweight container solution over virtual machines when identical programs are tested in both environments [11, 20].

This suggests that a container-based approach is likely preferable if the choice is available. However, not every application can be run in a container-based environment, so we maximize coverage of potential usecases by offering both container-based and virtual machine-based application hosting.

Regarding Platform-as-a-Service (PaaS), studies show that there is no dramatic difference in performance when derivatives of operating systems are compared to each other [1, 3]. Nevertheless, some derivatives are more lightweight and performant. In contrast, comparing operating systems of different types reveals more obvious results, particularly when considering low-level operating

¹ On Linux operating systems, sound and microphone devices are accessed via sound servers like PulseAudio and PipeWire, which interact with the sound interface present in the kernel, e.g., Advanced Linux Sound Architecture (ALSA) or Open Sound System (OSS)

system functionality implementations such as interrupt handling, memory allocation, and boot time [22]. Additional performance potentials can be leveraged in cases where a Windows virtual machine can be replaced with a container-based solution integrating wine [15].

For SaaS, performance largely depends on the underlying usage model and the actual software being virtualized. Depending on the use case, influential properties such as multi-tenancy, provision strategies, or the virtualized data type might be utilized to varying degrees, resulting in overall better or worse performing systems [25].

Despite the potential for usecase-specific optimization, we emphasize that for DaaS systems, coverage of a broad range of potential use cases is generally more favorable than solving a particular use case most efficiently. Therefore, we consider all user-facing desktop applications as plausible candidates for migration to our systems, rather than more complex distributed systems for which tailored environments might be more suitable.

Figure 2 visualizes the available hosting strategies for our platform. We aim to support any Windows or Linux-based application in virtual machines and any Linux- and Wine-based applications in container environments.

DaaS	DESIGN Frontend			
SaaS	Linux App		Windows App	
PaaS	Linux OS		Windows OS	Linux OS, Wine
IaaS	Virtual Machine	Container	Virtual Machine	Container

Fig. 2. Application hosting strategies.

Network Latency and Bandwidth Network Latency and Bandwidth are crucial factors in delivering web-based access to DaaS apps, directly impacting system responsiveness experienced by users. Supporting standardized frame-buffer protocols enables user interaction initially and provide additional access to external devices like printers or USB sticks. In cases involving, for example, video content transmission, network bandwidth often becomes the main bottleneck, especially with high-resolution video streams. In contrast to that, when transmitting audio, mouse movements, or keyboard inputs, network latency becomes more important.

Application-specific requirements can be managed to some extent through protocol-inherent tuning parameters and compression algorithms. Our solution, therefore, offers a range of tuning options to adapt to different use cases on a per-user basis.

We utilize the Guacamole desktop gateway implementation, supporting widely used frame-buffer protocols like VNC and RDP, along with relevant tuning options. Guacamole also enables audio support and integration with external devices through a JavaScript-based web client (guacamole.js). As depicted in figure 3, the proxy component receives HTTP requests from guacamole.js and is forwarded by the backend through WebSocket opcodes (Guacamole protocol). These opcodes are then transmitted to the Guacamole daemon (guacd) which in turn communicates to each instance by using regular frame-buffer protocols (VNC or RDP).

However, integrating Guacamole into the system requires consideration of additional factors besides the frame-buffer protocols itself. Authentication mechanisms, for example, are essential in a DaaS context, as users expect seamless authentication across the whole system without multiple logins. Additionally, safeguarding sensitive information like user passwords is crucial for maintaining system security and user trust. To mitigate the risk of unintended system manipulation, Guacamole was integrated using a proxied approach and WebSockets. This ensures that the credentials of a generated instance remain secret to the host system while still providing authorized, passwordless access to the hosted application.

Moreover, the proxy mechanism enables support for other additional features such as dynamic desktop resizing. While Guacamole offers support for that by default in RDP, support for the VNC protocol is completely missing. Nevertheless, the proxied approach enables intercepting all exchanged WebSocket opcodes and by listening for size requests the requested feature could be added by directly interacting with the instance on another layer.

While these changes to the default Guacamole use case are necessary to align with the underlying DaaS design requirements, they also negatively impact overall system performance and responsiveness. The integration as a proxy introduces additional network hops, leading to noticeable delays in network communication. Additionally, the protocol-level evaluation of opcodes further increases this delay, potentially affecting perceived responsiveness if resize requests take longer than expected. However, despite the introduced overhead, the solution still enables platform-independent and seamless access for all defined DaaS use cases.




Fig. 3. WebSocket Proxy Components.

Internal processing and Communication Another critical aspect directly related to performance and responsiveness is how data is internally processed, stored, and communicated. This includes metadata collected and processed and involves data exchange with all relevant backend services or instances.

Ensuring a satisfying user experience in a DaaS system implies specific components to be available. This includes access to all infrastructural backend services, a databases, authentication services, filesystems, and network devices. Additionally, a web server is integrated to provide relevant API endpoints for front-end communication and WebSocket endpoints for the viewer.

Generally spoken, all of these components can potentially become bottlenecks when handling high throughputs or large amounts of concurrent applications and is especially true if all components are served from the same hardware. Architecturally, all relevant components must therefore be designed for distribution across different hosts. However, keeping specific components directly on the host system may still be required to ensure low-latency access.

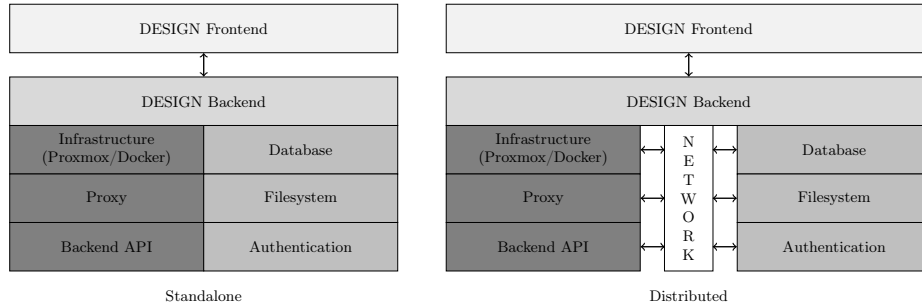
From our observations, components requiring real-time and low-latency performance, should physically be kept as close to the core hosting system as possible as shown in figure 4. This includes all data and communication relevant to audio, video, mouse, and keyboard transmission, as well as all backend services for hosting containers and virtual machines. Ideally, these components should be hosted natively on the host system for virtual network and hardware access or at least be connected to the same physical network to minimize network latencies. All other components might be distributed to dedicated systems if large user amounts and specific use cases have to be handled but keeping them in close physical proximity to the host remains prudent in smaller scales.

Although we can not yet provide a full-fledged performance analysis, we already measured promising results in our local test environment. Depending on the exact setup, the solution provides reasonable round-trip times and delays for each communication direction visualized in figure 4. In a minimal setup, we measured delays of around 10ms for client communication, below 1ms for internal backend communication but up to 1000ms for instance communication. But these numbers are only rough estimates as certain optimization concepts remain to be implemented.

4.3 Stability

Stability in DaaS applications is especially important for meeting the demands of continuous and time-critical operations across diverse user scenarios. Ensuring the stability of all system components during runtime is paramount to maintain service continuity, handle growing user loads, and transparent failure recovery while providing a seamless end-user experience. These considerations are integral to our architectural concept, directly impacting the operational integrity of the system and its suitability for enterprise-level deployment and trust.

Therefore, key strategies, such as component redundancy, node scalability as well as the internal communication concept will be described in the following subsections.

**Fig. 4.** Internal Communication Flow

Scalability of Components Scalability in DaaS environments can be approached through two fundamental methods: horizontal and vertical scaling. Horizontal scaling (scale out), involves adding more machines or instances to manage increased loads. In contrast, vertical scaling (scale up), adds more power (CPU, RAM) to existing machines [23]. While horizontal scaling offers flexibility and is often preferred in cloud environments, vertical scaling remains important for applications requiring strong single-thread performance or low latencies.

To comply with both approaches, we designed the system and its backend services to be split into logical units that can be distributed across dedicated systems. Each backend service receives a dedicated configuration and can be accessed via hostname and port tuples. By default, all backend services reside on the same host. In a distributed context, services can be distributed to dedicated remote systems to adhere to horizontal scaling principles.

Vertical scaling of DaaS nodes is not currently planned, but administrators are free to use custom third-party toolchains, such as GPU virtualization for vertical scaling on a per-host basis. Although vertical scaling within a host is beyond the scope of our research project, scaling the instances hosted within the system is feasible for both dimensions. For vertical scaling of instances, we provide appropriate parameters and API endpoints to manipulate resource assignments, allowing for arbitrary configurations.

For horizontal scaling of instances, the built-in mechanisms of each backend service are automatically configured and utilized to distribute instances across a pool of connected nodes. Proxmox implements such mechanisms for virtual machines as part of its ‘Proxmox VE High Availability (HA)’ cluster. Similar concepts can be integrated for containers with any backend service, such as Kubernetes. Within these clusters, DaaS instances are entirely agnostic to their hosting system and only need to be reachable via a known hostname or IP as a single requirement. As shown in figure 5, we provide a globally maintained state of data, authentication, and storage which is accessible either by horizontally scaled backends and frontends or otherwise, by horizontally and vertically scaled instances in our DaaS object domain.

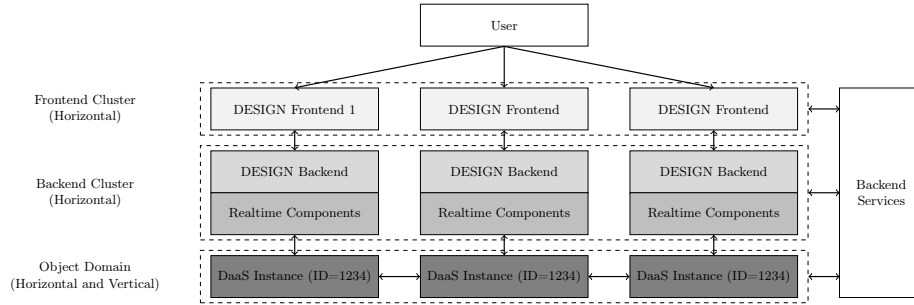


Fig. 5. DESIGN Component Scaling

Redundancy of Backend Services Redundancy is crucial in distributed systems, particularly in DaaS architectures, to safeguard against service interruptions and data loss by duplicating essential components, functions, and data. This design principle ensures high availability and reliability, even during component failures. Figure 6 depicts all major DESIGN components eligible for redundancy and which will be briefly described following:

- **Frontend/Backend:** In our final version, redundancy of the frontend and backend components will enable configuration for distributing multiple instances to dedicated systems and utilizing standardized load-balancing strategies. This setup allows the frontend or backend to serve as a gateway between the user and the backend service containing all relevant instances. Both components contribute to redundancy by distributing requests and responses to their respective recipients instead of processing them on their own.
- **Instances:** Redundancy of instances It is not a necessary condition for DaaS systems in general, but applying such strategies under certain conditions might be plausible. By using highly available clusters within our backend services redundancy can be achieved by using the corresponding API endpoints previously described and by parameterizing them appropriately.
- **Database:** The database is integrated as a backend service and is, therefore, independent of surrounding components and contexts and can be externalized to a dedicated systems. Modern database solutions, like MariaDB, offer rich features for maintaining specific data consistency strategies such as data replication, clustering, failover mechanisms or RAID control. Currently, redundancies in our database system are not planned but can be easily applied by using such database-specific features and a modified database adapter.
- **Authentication:** Authentication within our architecture is implemented as a backend service, also allowing for distribution to dedicated systems. It can be targeted by standardized load-balancing mechanisms and remains entirely agnostic of calling contexts. The system maintains its credential registry in a separate database, along with customizable user and group privileges and implements the well-defined OAuth 2.0 standard to facilitate remote authentication.

- **Distributed Filesystem:** In our architecture, we utilize the distributed filesystem Ceph, which offers means of object, block, and file storage. Redundancy is implicitly achieved through its self-healing and self-managing features, employing the CRUSH algorithm (Controlled Replication Under Scalable Hashing). Data availability and fault tolerance are ensured by automatically replicating data across multiple systems. Such benefits are leveraged by providing reliable file storage to all instances, and by providing block storage to the infrastructural backend services. By doing that, migration and redundancy of instances is achieved or, in other words, horizontal scaling.

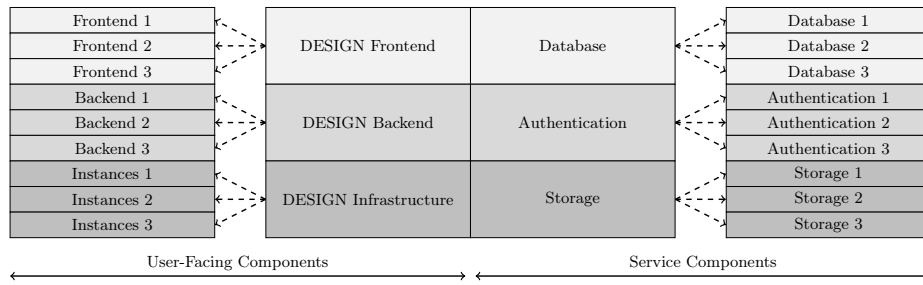


Fig. 6. DESIGN Component Redundancy

Internal Communication After discussing the principles of scalability and redundancy, we finally describe how communication flows are managed within our system. Our solution is rooted in a message-oriented architecture designed to fulfill DaaS properties for scalability and redundancy, as previously outlined.

This implies that any action triggered, whether in the frontend on the client side, in the backend on the host side, or within hosted instances, must ensure that all necessary information is readily available at the time of processing.

We achieve this by storing runtime information in distributed block storage, relevant configuration files in distributed file storage, object metadata in a common database, and by utilizing a decentralized authentication component.

This ensures that any change triggered by a node within the DaaS domain is immediately available to any other node relying in some way on that particular information change. Additionally, any extra information needed by a component must be fully contained within the related request. As shown in figure 7, this enables arbitrarily distributed configurations as long as request and response messages reach their desired target eventually.

If that is the case, each distributed host system might only contain its essential components, including the backend service with all hosted instances as well as the described proxy component for direct user interaction. Any other component may reside somewhere in the DaaS network.

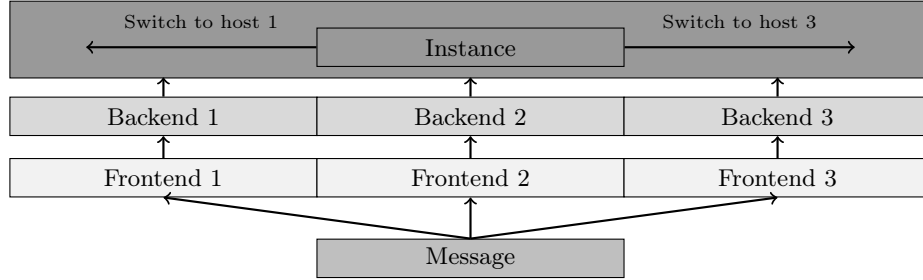


Fig. 7. Message Oriented Routing.

4.4 Usability

An essential feature of a DaaS platform is its strong user-friendliness, which eliminates the complexity of the different operating systems used inside virtual machines and containers, makes the underlying hardware transparent, and simplifies the necessary interactions as much as possible.

Since users already have an operating system on their clients, and the motivation for using a DaaS is the simple use of various applications, it makes no sense for most users to have additional desktops available in the browser. The focus should be on the applications. For this reason, when developing a new DaaS solution, the main goal should be to reduce the graphical output to the limits of the respective application. The concept uses a separate browser tab for each application whose content exclusively represents the application.

Developing the automatic adaptation of the output to the limits of the respective application window is challenging. Only a few RDP and VNC server implementations offer to export the graphical output of only individual windows or processes. For Windows operating systems, the proprietary built-in RDP server and the free software TightVNC offer the export of full desktops and single windows.

For Linux operating system deployments, the free software projects xrdp and x11vnc enable all desired features and single-app functionality can be easily obtained by using a custom window manager such, as for example, qtile ².

Table 1 includes an overview of the mentioned service implementations for RDP and VNC protocols and their relevant characteristics. Further implementations like TigerVNC and protocols like NX NoMachine exist but do not seem useful for designing and implementing a novel DaaS because they do not meet all feature requirements [2].

The automatic adaptation of viewer size and screen resolution, considering underlying protocols and systems, faced various limitations, briefly described below.

One challenge was imposed by Guacamole’s current lack of support for dynamic desktop resizing with VNC, leading to misaligned layouts and a poor user

² <https://qtile.org/>

Table 1. Services used by our DaaS for Remote Access to individual Applications

Service implementation	Protocol	Single Window Mode	Operating System	Software License
TightVNC	VNC	supported	Windows	GPL2
Windows RDP server	RDP	supported	Windows	proprietary
x11vnc	VNC	not supported	Linux	GPL2
Xrdp	RDP	supported	Linux	Apache 2.0

experience. To address this until such features are available in Guacamole ³, we implemented a separate mechanism using a WebSocket proxy and direct interaction with each instance.

Additionally, protocol-specific behaviors may vary when hosted on different platforms or infrastructures. For instance, we observed connection terminations when resizing Linux-based containers, contrasting with the behavior observed in equivalent Linux-based virtual machines. To overcome such limitations, we created a separate layer around the Guacamole client, adapting to each infrastructure, platform, or protocol-specific configuration.

Another challenge was managing RDP sessions, which are typically tied to actual users, unlike VNC sessions, which are typically authenticated against a server-maintained user registry. Utilizing the WebSocket proxy, the backend could pass credentials to the server without revealing them to the user.

Essentially, Guacamole, with its protocol implementations and additional features like audio and printing, provided the fundamental base for our core capabilities.

Finally, certain instances might not be fully accessible during runtime due to hardware malfunctions and inappropriate system configurations during system updates. Therefore, as a failover mechanism, additional logic had to be added to provide reasonable debug views and recovery mechanisms, e.g., recreate instances from a specific snapshot or configuration state.

A crucial lesson learned was, therefore, to realize that while individual problems can mostly be solved using established paradigms, services, or libraries with only slight modifications, the real challenge for a full-fledged DaaS system lies in its abilities to combine such isolated solutions appropriately and for each specific infrastructure and platform configuration.

5 Conclusion and Outlook

In this paper, we describe an architecture for a novel and feature-rich DaaS that is, by principle, superior to comparable products and approaches because it supports running unmodified Linux and Windows applications. Furthermore, we identified characteristic issues – concerning compatibility, performance, stability, and usability – during the development and implementation of our DaaS and present solutions to overcome these challenges.

³ <https://issues.apache.org/jira/browse/GUACAMOLE-1196>

For future work, we plan to replace our temporary solution for internal communication with a more efficient approach as our preliminary performance tests revealed a potential bottleneck in that particular area. We also plan to apply certain optimization strategies with respect to backend and instance communication. We further plan to enhance our yet limited performance testing and will provide a full-fledged performance analysis of realistic deployment scenarios and utilization patterns which include an evaluation of different user behaviors within various popular applications.

Acknowledgements

The Federal Ministry for Economic Affairs and Climate Action funded this work (*'Bundesministerium für Wirtschaft und Klimaschutz'*) in the framework of the central innovation programme for small and medium-sized enterprises (*'Zentrales Innovationsprogramm Mittelstand'*).

We thank our project partners from Nuromedia GmbH for their support. We especially thank Björn Goetschke, Mike Ludemann, Dario Savella, Holger Sprengel, and Rahul Tomar.

References

1. Balen, J., Vdovjak, K., Martinović, G.: Performance evaluation of windows virtual machines on a linux host. *Automatika: časopis za automatiku, mjerenje, elektroniku, računarstvo i komunikacije* **61**(3), 425–435 (2020)
2. Baun, C., Bouche, J.: Closing the gap between web applications and desktop applications by designing a novel desktop-as-a-service (daas) with seamless support for desktop applications. *Open Journal of Cloud Computing (OJCC)* **8**(1), 1–19 (2023)
3. Boras, M., Balen, J., Vdovjak, K.: Performance evaluation of linux operating systems. In: *2020 International Conference on Smart Systems and Technologies (SST)*. pp. 115–120. IEEE (2020)
4. Carrascosa, M., Bellalta, B.: Cloud-gaming: Analysis of google stadia traffic. *Computer Communications* **188**, 99–116 (2022)
5. Celesti, A., Mulfari, D., Fazio, M., Villari, M., Puliafito, A.: Improving desktop as a service in openstack. In: *2016 IEEE Symposium on Computers and Communication (ISCC)*. pp. 281–288. IEEE (2016)
6. Choy, S., Wong, B., Simon, G., Rosenberg, C.: The brewing storm in cloud gaming: A measurement study on cloud to end-user latency. In: *2012 11th Annual Workshop on Network and Systems Support for Games (NetGames)*. pp. 1–6 (2012). <https://doi.org/10.1109/NetGames.2012.6404024>
7. Claypool, M., Claypool, K.: Latency can kill: precision and deadline in online games. In: *Proceedings of the first annual ACM SIGMM conference on Multimedia systems*. pp. 215–222 (2010)
8. Clincy, V., Wilgor, B.: Subjective evaluation of latency and packet loss in a cloud-based game. In: *2013 10th International Conference on Information Technology: New Generations*. pp. 473–476. IEEE (2013)

9. Darcy, R.: Latency and Jitter in Cloud Game Streaming. Ph.D. thesis, Worcester Polytechnic Institute (2022)
10. Di Domenico, A., Perna, G., Trevisan, M., Vassio, L., Giordano, D.: A network analysis on cloud gaming: Stadia, geforce now and psnow. *Network* **1**(3), 247–260 (2021)
11. Felter, W., Ferreira, A., Rajamony, R., Rubio, J.: An updated performance comparison of virtual machines and linux containers. In: 2015 IEEE international symposium on performance analysis of systems and software (ISPASS). pp. 171–172. IEEE (2015)
12. Garmpis, A., Gouvatsos, N.: Design and development of webubu: An innovating web-based instruction tool for linux os courses. *Computer Applications in Engineering Education* **24**(2), 313–319 (2016)
13. Guacamole, A.: Configuring guacamole (2024), <https://guacamole.apache.org/doc/1.5.5/gug/configuring-guacamole.html>
14. Han, B.: Implementation of a Payout Buffer for a Cloud-based Game Client. Ph.D. thesis, Worcester Polytechnic Institute
15. Huang, C., Chen, J., Zhang, L., Luo, Q.: Performance evaluation of virtualization technologies for windows programs running on linux operating system. In: International Conference on Network Computing and Information Security. pp. 759–766. Springer (2012)
16. Jarschel, M., Schlosser, D., Scheuring, S., Hofkfeld, T.: An evaluation of qoe in cloud gaming based on subjective tests. In: 2011 Fifth International Conference on Innovative Mobile and Internet Services in Ubiquitous Computing. pp. 330–335. IEEE (2011)
17. Lampe, U., Wu, Q., Dargutev, S., Hans, R., Miede, A., Steinmetz, R.: Assessing latency in cloud gaming. In: Cloud Computing and Services Science: Third International Conference, CLOSER 2013, Aachen, Germany, May 8–10, 2013, Revised Selected Papers 3. pp. 52–68. Springer (2014)
18. Liu, K., Dong, L.j.: Research on cloud data storage technology and its architecture implementation. *Procedia Engineering* **29**, 133–137 (2012)
19. Magana, E., Sesma, I., Morato, D., Izal, M.: Remote access protocols for desktop-as-a-service solutions. *PloS one* **14**(1), e0207512 (2019)
20. Potdar, A.M., Narayan, D., Kengond, S., Mulla, M.M.: Performance evaluation of docker container and virtual machine. *Procedia Computer Science* **171**, 1419–1428 (2020)
21. Sefraoui, O., Aissaoui, M., Eleuldj, M., et al.: Openstack: toward an open-source solution for cloud computing. *International Journal of Computer Applications* **55**(3), 38–42 (2012)
22. Sulaiman, N.S., Raffi, A.S.H.A.: Comparison of operating system performance between windows 10 and linux mint. *International Journal of Synergy in Engineering and Technology* **2**(1), 92–102 (2021)
23. Vaquero, L.M., Roderio-Merino, L., Buyya, R.: Dynamically scaling applications in the cloud. *ACM SIGCOMM Computer Communication Review* **41**(1), 45–52 (2011)
24. VidyaBanu, R., Preethi, J., Dinesh, N.: Implementation of financial system using EyeOS in the cloud environment. In: 2011 International Conference on Recent Trends in Information Technology (ICRTIT). pp. 656–660. IEEE (2011)
25. Zhong, L., Wo, T., Li, J., Li, B.: A virtualization-based saas enabling architecture for cloud computing. In: 2010 Sixth International Conference on Autonomic and Autonomous Systems. pp. 144–149. IEEE (2010)